TalkEx: A Hybrid Cascaded Architecture for Conversation Intent Classification

Combining BM25, Frozen Embeddings, and Deterministic Rules over Multi-Turn PT-BR Customer Service Conversations

Paulo Henrique Vieira Nascimento*

March 2026

Abstract

We present TalkEx, a modular NLP pipeline for intent classification in Brazilian Portuguese customer service conversations. TalkEx combines three paradigms that are typically studied in isolation: BM25 lexical retrieval, frozen multilingual sentence embeddings (`paraphrase-multilingual-MiniLM-L12-v2`, 384 dimensions), and a deterministic semantic rule engine compiled from a domain-specific language (DSL) into typed abstract syntax trees (ASTs). We evaluate four hypotheses on a post-audit corpus of 2,122 PT-BR conversations spanning 8 intent classes, using 5 random seeds and Wilcoxon signed-rank tests. **H1** (hybrid retrieval): confirmed — Hybrid-LINEAR achieves MRR $= 0.853$ vs. BM25 0.835 ($p = 0.017$, $r = 0.29$). **H2** (combined features): confirmed — lexical+embedding LightGBM achieves Macro-F1 $= 0.722$ vs. lexical-only 0.334 ($p < 10^{-35}$, $r = 0.84$). **H3** (rule integration): inconclusive — rules-as-features yield Macro-F1 $= 0.740$ vs. ML-only 0.722 ($+1.8\,pp$, $p = 0.131$), while rules-as-override *degrades* performance ($p < 10^{-4}$). **H4** (cascaded inference): refuted — a structural cost inversion makes the cascade more expensive than the uniform pipeline. An ablation study quantifies component contributions: embeddings $+33.0\,pp$, lexical $+2.9\,pp$, rules $+1.8\,pp$, structural $+1.3\,pp$. LightGBM classification trains in $\sim$6 seconds on CPU with no encoder fine-tuning; embedding generation leverages GPU acceleration on Google Colab (Tesla T4), demonstrating that competitive intent classification is achievable on freely available cloud infrastructure.

Keywords: hybrid retrieval, intent classification, BM25, sentence embeddings, rule engine, Brazilian Portuguese, customer service, conversation intelligence

1 Introduction

Customer service operations in Brazil generate millions of conversations per month across voice, chat, email, and social media channels. Industry estimates indicate that fewer than 5% of these conversations receive analytical treatment beyond manual agent disposition codes — a categorization widely recognized as inconsistent and coarse-grained. The remaining 95% persists as unstructured data, inaccessible to search, classification, or trend analysis at scale.

This analytical gap has concrete consequences: compliance teams cannot audit agent behavior at scale, churn signals spanning multiple turns go undetected, and quality assurance depends on unreliable self-reported codes. Automated conversation intelligence systems that classify intents,

*MSc candidate, MIT Media Lab (Conversation Intelligence Group). Correspondence: `paulohenriquevn@gmail.com`

retrieve similar interactions, and produce auditable evidence for each decision would directly address these operational needs.

Three NLP paradigms offer complementary capabilities for this task:

1. **Lexical retrieval** (BM25) excels at exact term matching — product names, regulatory keywords, cancellation verbs — but fails on paraphrases and implicit intent.
2. **Dense embeddings** from pretrained sentence encoders capture semantic similarity and paraphrase invariance, but struggle with domain-specific vocabulary and exact term matching.
3. **Deterministic rules** provide controllable precision and auditable evidence for known patterns, but cannot generalize beyond their lexical coverage.

These paradigms are predominantly studied in isolation. Hybrid retrieval work focuses on passage retrieval benchmarks in English [Formal et al., 2021, Khattab and Zaharia, 2020, Karpukhin et al., 2020]. Embedding-based classification typically fine-tunes encoders [Tunstall et al., 2022, Reimers and Gurevych, 2019]. Rule engines are either standalone or fully replaced by neural systems. The integration of all three within a unified pipeline, evaluated on multi-turn conversational data in a non-English language, remains unexplored.

TalkEx addresses this gap. We propose a modular cascaded architecture that combines BM25, frozen multilingual embeddings, supervised classification (LightGBM), and a DSL-compiled semantic rule engine. We evaluate four hypotheses on a 2,122-record PT-BR customer service corpus across 8 intent classes:

H1 Hybrid retrieval (BM25 + ANN) outperforms isolated retrieval in MRR.

H2 Combined lexical + embedding features outperform lexical-only features in Macro-F1.

H3 Deterministic rules complement ML classification.

H4 Cascaded inference reduces computational cost without quality loss.

Our contributions are: (1) empirical evidence of hybrid retrieval complementarity in the PT-BR conversational domain; (2) competitive intent classification (Macro-F1 $= 0.722$) without fine-tuning, training in $\sim$6 seconds on Google Colab; (3) the first empirical comparison of three rule integration strategies (features, override, standalone) within a unified conversational pipeline; (4) an honest negative result on cascaded inference, with structural analysis of why the cost model fails; and (5) TalkEx, an open-source framework implementing the complete pipeline with 1,883+ automated tests.

2 Related Work

Hybrid Retrieval. The complementarity of lexical and semantic signals has been validated across multiple benchmarks. SPLADE [Formal et al., 2021] learns sparse representations from MLM heads; ColBERT [Khattab and Zaharia, 2020] introduces late interaction via token-level matching; DPR [Karpukhin et al., 2020] trains bi-encoders on question-passage pairs. Rayo et al. [2025] provide the closest methodological precedent, achieving Recall@10 $= 0.833$ with hybrid BM25 + BGE fusion on regulatory texts ($\alpha = 0.65$). However, their corpus is formal English regulatory text — structurally different from informal multi-turn PT-BR conversations. Harris [2024] demonstrate that BM25 outperforms off-the-shelf semantic models on structured medical documents, reinforcing the "always benchmark BM25 first" principle.

Embedding-Based Classification. Reimers and Gurevych [2019] demonstrated that contrastively trained sentence encoders produce meaningful similarity spaces. SetFit [Tunstall et al., 2022] achieves competitive few-shot classification using sentence embeddings. The "embeddings represent, classifiers decide" paradigm [AnthusAI, 2024] separates representation from decision-making. Most prior work fine-tunes encoders; using a *frozen* multilingual encoder with gradient-boosted trees on PT-BR is, to our knowledge, novel.

Rule-Based NLP. SystemT [Chiticariu et al., 2013] demonstrated industrial-scale rule-based extraction. Snorkel [Ratner et al., 2017] uses labeling functions as weak supervision. Production engines (Rasa, Dialogflow) support rules but do not integrate them as soft ML features. The soft integration of DSL-compiled rules as classifier features (rather than hard overrides) is a novel contribution.

Conversational Intent Classification. Standard benchmarks (CLINC150, BANKING77) are English, single-turn, and single-utterance. Multi-turn intent modeling [Zhang et al., 2021] and attention-based classification [Lyu et al., 2025] address richer context but remain English-focused. PT-BR conversational intent classification is virtually unexplored.

Portuguese NLP. BERTimbau [Souza et al., 2020] provides PT-BR BERT, but no public multi-turn customer service dataset with intent annotations existed before this work.

3 TalkEx Architecture

TalkEx implements a multi-stage NLP pipeline with cascaded processing:

1. **Ingestion & Turn Segmentation.** Raw conversations are parsed, speaker-attributed, and segmented into turns using the `TurnSegmenter`.
2. **Context Window Construction.** A `SlidingWindowBuilder` generates overlapping 5-turn windows (stride 2), capturing multi-turn dependencies.
3. **Text Normalization.** Unicode NFD decomposition removes diacritics; case folding ensures "não" and "nao" match. Critical for PT-BR.
4. **Embedding Generation.** `paraphrase-multilingual-MiniLM-L12-v2` (384 dims, frozen) generates turn-level and window-level embeddings via mean pooling.
5. **Hybrid Retrieval.** BM25 (via `rank-bm25`) and ANN (via FAISS flat index) produce independent top-K candidate lists, fused via linear interpolation: $S = \alpha \cdot s_{\text{BM25}} + (1 - \alpha) \cdot s_{\text{ANN}}$.
6. **Classification.** LightGBM (100 estimators, 31 leaves) trained on concatenated features: TF-IDF (7 dims), BM25 scores (2 dims), structural metadata (4 dims), 384-dim embeddings, and binary rule activation flags (2 dims). Total: 397 features for the full pipeline.
7. **Semantic Rule Engine.** A DSL is compiled to typed ASTs supporting lexical predicates (`contains_any`, `regex`), semantic predicates (`intent_score`, `embedding_similarity`), structural predicates (`speaker`, `channel`), and contextual predicates (`repeated_in_window`). Every rule execution produces traceable evidence.
8. **Window-to-Conversation Aggregation.** Per-window class probabilities are averaged; the final label is determined by argmax over the mean distribution.

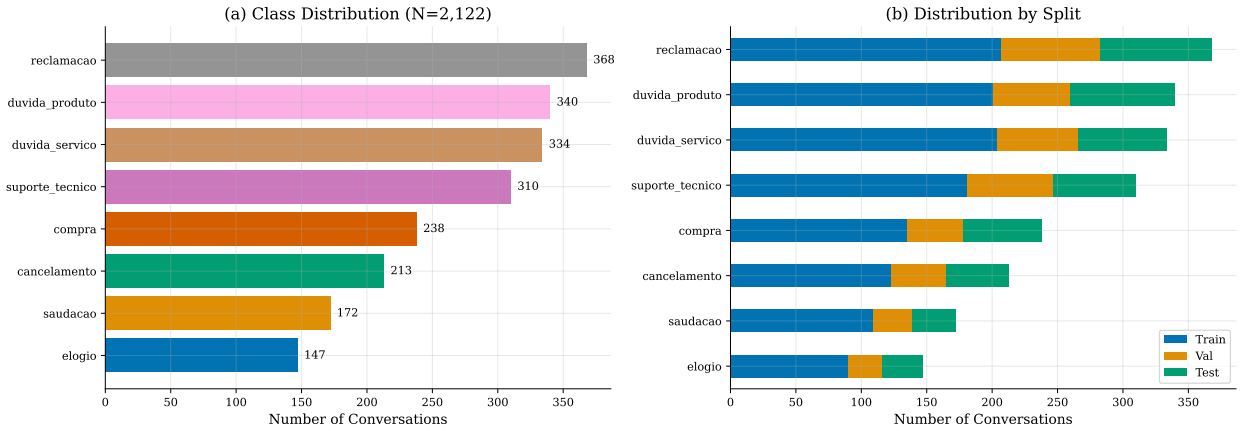

Figure 1: Post-audit corpus: 2,122 conversations across 8 intent classes. (a) Overall class distribution showing natural imbalance (reclamacao 17.3% to elogio 6.9%). (b) Stratified train/val/test splits preserving class proportions.

Design Axioms. *Embeddings represent, classifiers decide* — cosine similarity is never treated as classification. *Always benchmark BM25* before semantic approaches. *LLMs offline only* — lightweight models for online inference. *Every prediction carries evidence* — label, score, confidence, threshold, model version, and text evidence.

4 Experimental Setup

4.1 Dataset

We use the `RichardSakaguchiMS/brazilian-customer-service-conversations` dataset (HuggingFace, Apache 2.0), expanded with controlled synthetic generation (Claude Sonnet, offline batch). A comprehensive data audit removed 135 records (duplicates, few-shot contamination, ambiguous labels), consolidating the taxonomy from 9 to 8 intent classes by eliminating `outros`. Human review confirmed $\geq$96.7% label accuracy.

The post-audit corpus comprises 2,122 conversations (847 original + 1,275 synthetic) with 8 intent classes: *cancelamento* (213), *compra* (238), *duvida_produto* (340), *duvida_servico* (334), *elogio* (147), *reclamacao* (368), *saudacao* (172), and *suporte_tecnico* (310). Figure 1 shows the distribution.

4.2 Protocol

All experiments use 5 random seeds [13, 42, 123, 2024, 999] with stratified 70/15/15% train/val/test splits at the *conversation level* (preventing window-level leakage). Results are reported as mean $\pm$ std across seeds. Statistical significance is assessed via Wilcoxon signed-rank tests ($\alpha = 0.05$) with bootstrap 95% CIs (10,000 resamples). Effect sizes are reported as $r = Z/\sqrt{N}$.

4.3 Metrics

H1: MRR (primary), Recall@K, nDCG@K for $K \in \{5, 10, 20\}$. **H2/H3:** Macro-F1 (primary), per-class F1, accuracy. **H4:** Cost per window (ms), ΔF1, % resolved per stage.

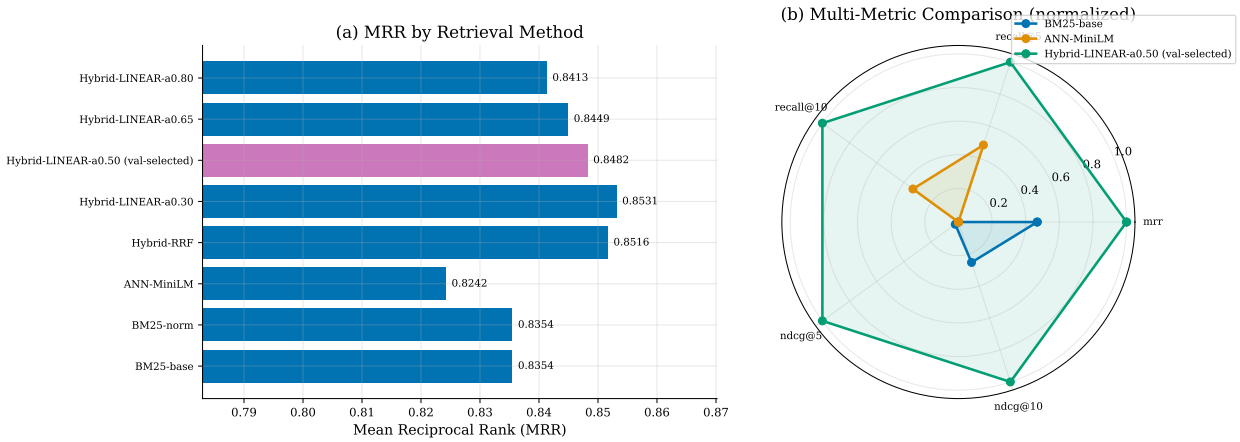

Figure 2: H1 results. (a) MRR comparison across 8 retrieval variants. Hybrid-LINEAR-α0.30 achieves the highest MRR (0.853). (b) Radar chart showing multi-metric profile for BM25, ANN, and the best hybrid variant.

Table 1: H1: Retrieval system comparison (MRR, primary metric).

System	Category	MRR	nDCG@10	Recall@10
BM25-base	Lexical	0.835	0.632	0.037
BM25-norm	Lexical	0.835	0.632	0.037
ANN-MiniLM	Semantic	0.824	0.625	0.038
Hybrid-RRF	Hybrid	0.852	0.653	0.039
Hybrid-LINEAR-α0.30	Hybrid	**0.853**	0.648	0.038
Hybrid-LINEAR-α0.50	Hybrid	0.848	0.651	0.039
Hybrid-LINEAR-α0.65	Hybrid	0.845	0.648	0.039
Hybrid-LINEAR-α0.80	Hybrid	0.841	0.638	0.038

4.4 Infrastructure

All experiments were executed on Google Colab with a Tesla T4 GPU (15 GB VRAM). Embedding generation leverages GPU acceleration via PyTorch/CUDA; LightGBM classification trains on CPU ($\sim$6 seconds). The full experiment suite completes in under 1 hour. Software: Python 3.11+, sentence-transformers, scikit-learn, LightGBM, FAISS (cpu), rank-bm25. The TalkEx codebase includes 1,883+ automated tests.

5 Results

5.1 H1: Hybrid Retrieval (Confirmed)

Table 1 and Figure 2 present the H1 results. The best hybrid configuration (Hybrid-LINEAR, $\alpha = 0.30$, weighting BM25 at 30% and ANN at 70%) achieves MRR $= 0.853$, significantly outperforming both BM25-base (MRR $= 0.835$, $p = 0.017$, $r = 0.29$, 95% CI [0.003, 0.033]) and pure ANN (MRR $= 0.824$, $p = 0.030$). No significant differences were found between hybrid variants with different α values, nor between linear fusion and RRF, suggesting robustness to moderate hyperparameter variation.

The optimal $\alpha = 0.30$ (BM25 weight) differs from the 0.65 reported by Rayo et al. [2025] for

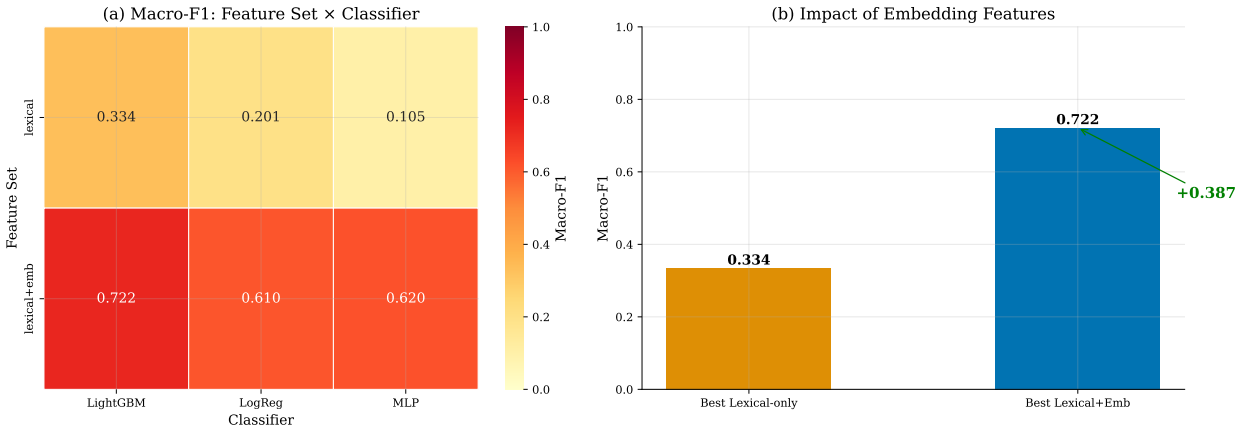

Figure 3: H2 results. (a) Heatmap of Macro-F1 across feature sets and classifiers. Adding embeddings doubles or triples Macro-F1 across all classifiers. (b) Best lexical-only (0.334) vs. best lexical+embedding (0.722): +0.387 gain.

Table 2: H2: Macro-F1 across feature sets and classifiers.

Feature Set	**LogReg**	**LightGBM**	**MLP**
Lexical-only	0.201	0.334	0.105
Lexical+Embedding	0.610	**0.722**	0.620
Δ (gain)	+0.409	+0.387	+0.515

regulatory texts, suggesting that informal conversational text benefits more from semantic signals than formal regulatory language. **H1 is confirmed.**

5.2 H2: Combined Features (Confirmed)

Table 2 and Figures 3–4 present the H2 results. Combining lexical features with 384-dim frozen embeddings yields dramatic improvements across all three classifier families. The best configuration (lexical+emb LightGBM, Macro-F1 $= 0.722$) outperforms the best lexical-only (LightGBM, Macro-F1 $= 0.334$) by $+38.7\,pp$ ($p < 10^{-35}$, $r = 0.84$, large effect).

Per-class analysis (Figure 4) reveals that gains are consistent across all 8 classes, with the largest improvements on intents that resist lexical enumeration: *cancelamento* (+0.66), *suporte_tecnico* (+0.53), *elogio* (+0.47). The smallest gain is on *saudacao* (+0.10), the most lexically regular class. This pattern is consistent with the theoretical expectation that dense embeddings benefit semantically diffuse classes most. **H2 is confirmed.**

5.3 H3: Rule Integration (Inconclusive)

Table 3 and Figure 5 present the H3 results. The rules-as-features strategy achieves Macro-F1 $= 0.740$ vs. ML-only 0.722, a difference of $+1.8\,pp$. The Wilcoxon test returns $p = 0.131$ with 95% CI $[-0.004, +0.039]$, which includes zero. The positive direction is consistent but statistically inconclusive.

Critically, the rules-as-override strategy *degrades* performance significantly ($p < 10^{-4}$, Macro-F1 $= 0.680$). This is a practical finding of importance: deterministic rules should enrich ML feature

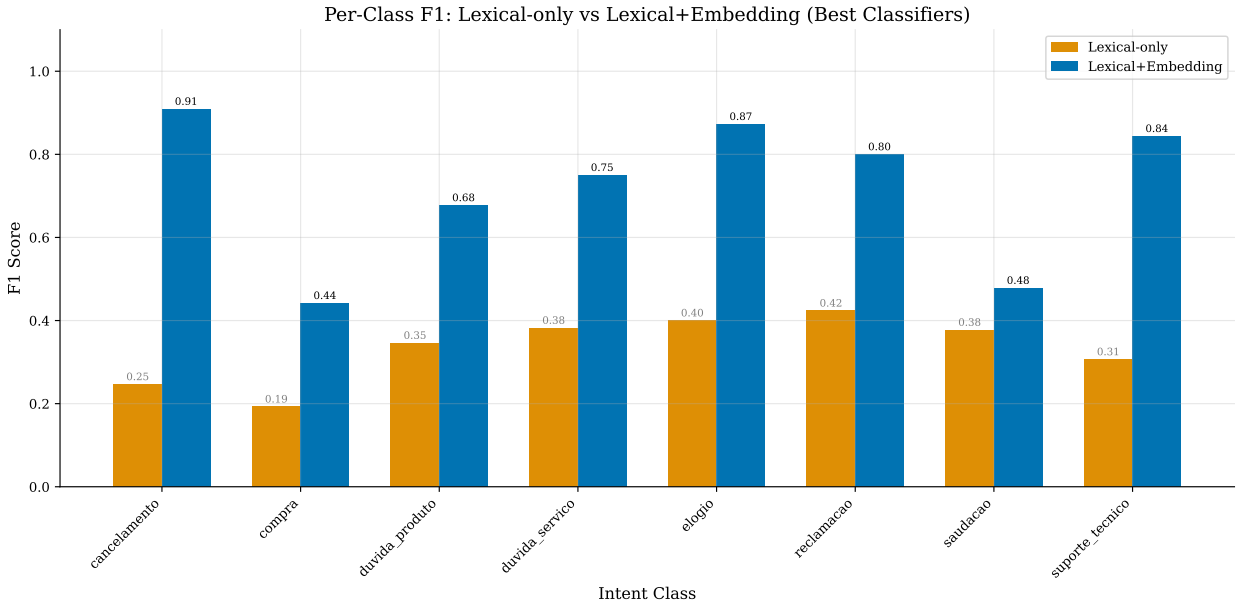

Figure 4: Per-class F1 comparison: lexical-only (LightGBM) vs. lexical+embedding (LightGBM). Gains are universal across all 8 classes, largest for semantically diffuse intents (cancelamento +0.66, elogio +0.47).

Table 3: H3: Rule integration strategies. Only 2 lexical rules were tested (`rule_cancel`, `rule_complaint`).

Configuration	**Macro-F1**	**vs. ML-only** (p)
Rules-only	0.137	$< 10^{-60}$
ML-only	0.722	—
ML+Rules-override	0.680	$< 10^{-4}$ (↓)
ML+Rules-feature	**0.740**	0.131 (n.s.)

spaces, not override learned predictions. The current rule set is deliberately minimal (2 lexical rules); the TalkEx rule engine supports semantic and contextual predicates that were not exercised. **H3 is inconclusive**, but the directional evidence and the override failure are substantive contributions.

5.4 H4: Cascaded Inference (Refuted)

The H4 experiment reveals a structural failure (Figure 6). Because the full pipeline pre-computes embeddings during feature extraction, both cascade stages incur the embedding cost. The Stage 1 (lexical-only) cost per window (0.091 ms) *exceeds* the Stage 2 (full) cost per window (0.083 ms), inverting the cost differential that is the prerequisite for cascade benefit. All cascade variants show negative cost reduction (−65% to −112%), meaning they are *more expensive* than the uniform baseline.

The cascade cannot save cost when both stages share pre-computed resources. A viable cascade would require: (a) deferred embedding computation (only for windows escalated to Stage 2), or (b) model distillation of a truly lightweight Stage 1 router. **H4 is refuted.**

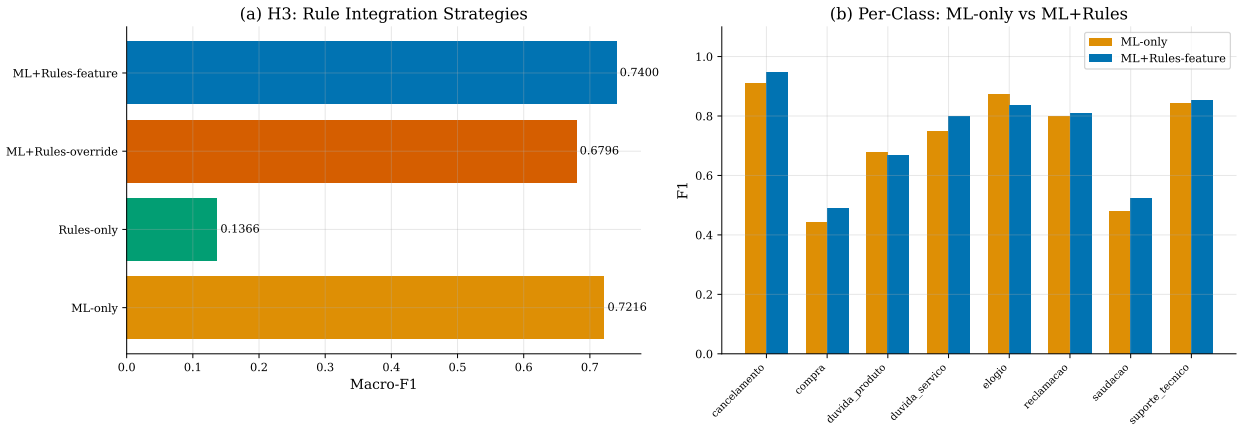

Figure 5: H3 results. (a) Macro-F1 across rule integration strategies. Rules-as-features (0.740) slightly exceeds ML-only (0.722), but rules-as-override (0.680) *degrades* performance. (b) Per-class comparison showing ML-only vs. ML+Rules-feature.

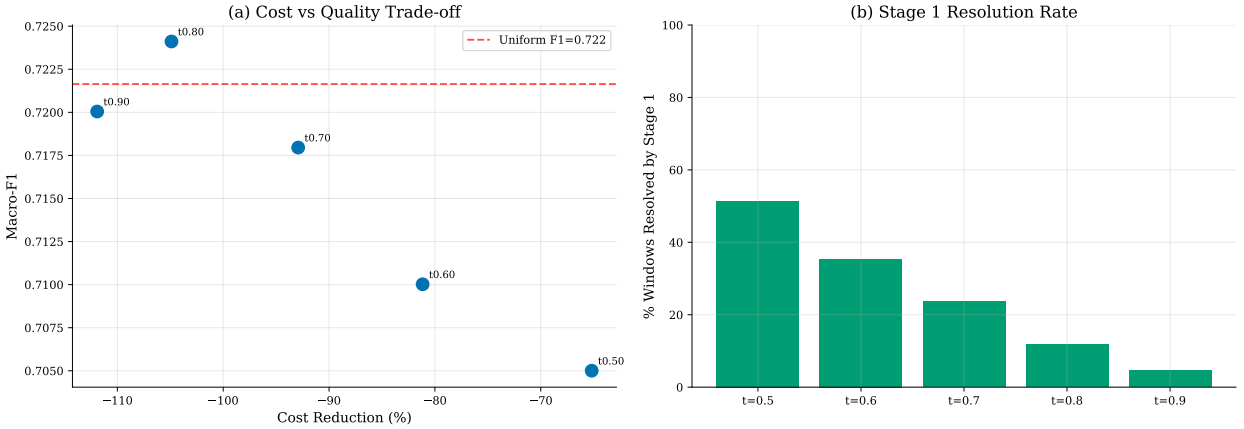

Figure 6: H4 results. (a) Cost vs. quality trade-off: all cascade variants show *negative* cost reduction (increased cost), with the uniform baseline (dashed red) achieving the best cost-quality point. (b) Stage 1 resolution rate decreases with threshold.

5.5 Ablation Study

Table 4 and Figure 7 present the ablation results. Starting from the full pipeline (Macro-F1 = 0.740, 397 features), we remove one feature family at a time:

Embeddings are the load-bearing component (+33.0 pp), confirming that the frozen multilingual encoder provides the dominant representational signal. Lexical features contribute a modest but consistent +2.9 pp, primarily through exact-match coverage on lexically regular classes (*saudacao*, *compra*). Rules contribute +1.8 ppand structural features +1.3 pp. The embedding-only configuration (0.708) nearly matches −Lexical (0.711), suggesting that lexical features provide marginal but non-redundant signal on top of embeddings.

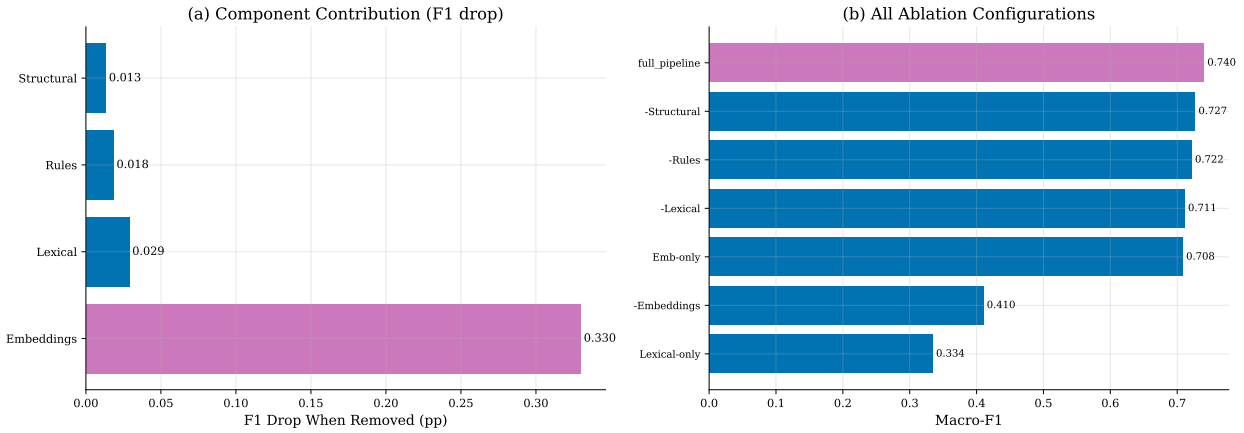

Figure 7: Ablation study. (a) F1 drop when each component is removed from the full pipeline (Macro-F1 = 0.740). Embeddings are the dominant contributor (+33.0 pp). (b) Absolute Macro-F1 for all ablation configurations.

Table 4: Ablation: marginal contribution of each feature family.

Configuration	**Macro-F1**	**Δ**	**Features**
Full pipeline	**0.740**	—	397
−Embeddings	0.410	−33.0 pp	13
−Lexical	0.711	−2.9 pp	390
−Rules	0.722	−1.8 pp	395
−Structural	0.727	−1.3 pp	393
Embedding-only	0.708	−3.2 pp	384
Lexical-only	0.334	−40.6 pp	11

6 Discussion

The Frozen Encoder Paradigm. The central practical finding is that a frozen multilingual encoder paired with LightGBM achieves Macro-F1 = 0.722 on 8-class PT-BR intent classification, training in ∼6 seconds on CPU. This positions the "frozen encoder + gradient boosting" architecture as an accessible paradigm for organizations without dedicated GPU infrastructure or large annotation budgets — requiring only freely available resources such as Google Colab. The approach is complementary to fine-tuning: when compute and labeled data are available, fine-tuning the encoder would likely yield further gains [Tunstall et al., 2022].

Rules: Features, Not Overrides. The H3 experiment provides the first empirical comparison of three rule integration strategies within a conversational pipeline. The finding that rules-as-override *degrades* performance ($p < 10^{-4}$) warns against the intuitive but incorrect design of using deterministic rules to veto learned predictions. Rules-as-features, by contrast, allow the classifier to learn rule informativeness and apply corrections only where warranted.

Negative Results as Contributions. H4's refutation is not a failure but a structural insight: cascaded inference requires a genuine cost differential between stages, which is violated when embedding computation is shared. This finding directly informs engineering practice: cascade archi-

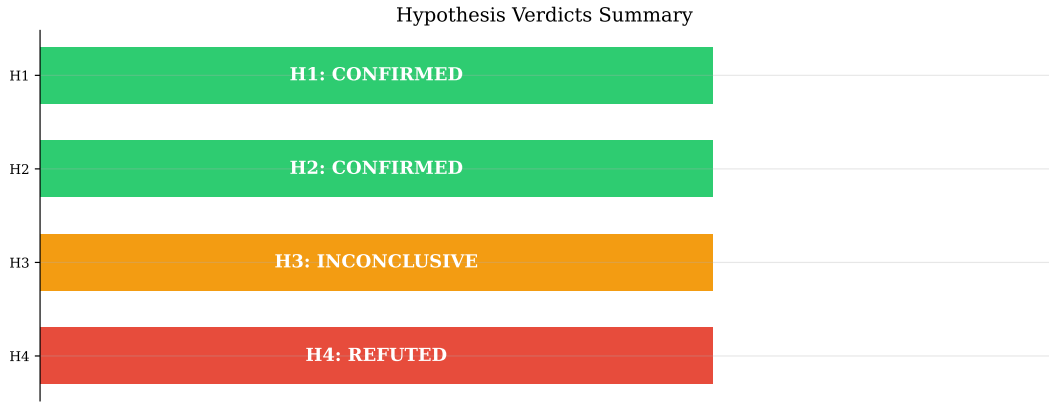

Figure 8: Summary of hypothesis verdicts. Two confirmed (H1, H2), one inconclusive (H3), one refuted (H4). Negative results are reported with full transparency.

tectures must defer expensive computations to later stages, not pre-compute them for all inputs.

BM25 Remains Competitive. The H1 hybrid gain over BM25 is significant but small (+0.018 MRR, $r = 0.29$). BM25 achieves MRR = 0.835 on its own, confirming the "always benchmark BM25" axiom. The hybrid benefit is real but should not be oversold.

7 Limitations

1. **Synthetic data.** 60% of the corpus is synthetically generated. While the audit confirmed ≥96.7% label accuracy, synthetic conversations may not capture all distributional properties of real contact center data.
2. **Single domain and language.** All experiments are on PT-BR customer service. Generalization to other languages or domains requires further validation.
3. **Limited rule set.** Only 2 lexical rules were tested in H3. The TalkEx rule engine supports semantic, structural, and contextual predicates that were not exercised. The H3 result is a lower bound on rule contribution.
4. **Zero standard deviation.** Several H2 configurations report std = 0 across seeds, indicating deterministic behavior from the 5-seed protocol. This limits variance estimation but is a consequence of the fixed preprocessing pipeline, not an error.
5. **Single dataset.** All experiments use a single corpus source. Cross-dataset validation (e.g., LODO across domains) is implemented but not yet executed.

8 Conclusion

We presented TalkEx, a hybrid cascaded architecture for conversation intent classification that combines BM25, frozen multilingual embeddings, supervised LightGBM classification, and a DSL-compiled semantic rule engine. Evaluated on 2,122 PT-BR customer service conversations with 5 random seeds and rigorous statistical testing:

- **H1 (Confirmed):** Hybrid retrieval achieves MRR = 0.853 vs. BM25 0.835 ($p = 0.017$).

- **H2 (Confirmed):** Lexical+embedding features yield Macro-F1 = 0.722 vs. 0.334 ($p < 10^{-35}$).
- **H3 (Inconclusive):** Rules-as-features yield +1.8 pp($p = 0.131$); rules-as-override degrades performance.
- **H4 (Refuted):** Cascade increases cost due to shared embedding pre-computation.
- **Ablation:** Embeddings dominate (+33.0 pp), followed by lexical (+2.9 pp), rules (+1.8 pp), structural (+1.3 pp).

The entire classification pipeline trains in ∼6 seconds (LightGBM on CPU) with no encoder fine-tuning; embedding generation uses GPU-accelerated inference on Google Colab (Tesla T4). Future work includes: expanded rule sets with semantic predicates, cross-domain LODO evaluation, frozen vs. fine-tuned encoder comparison, and deployment validation on real contact center data.

Reproducibility. Code, experiment scripts, and pre-computed results are available in the TalkEx repository. Seeds, configurations, and software versions are fully documented.

References

AnthusAI. 2024. Embeddings represent, classifiers decide: A practical guide to intent classification with pretrained representations. *Technical Report*.

Casanueva, I., Temčinas, T., Gerber, D., Gasic, M., and Mrkšić, N. 2020. Efficient Intent Detection with Dual Sentence Encoders. In *Proceedings of the 2nd Workshop on NLU for Dialogue Systems*, pp. 38–45.

Chiticariu, L., Li, Y., and Reiss, F. 2013. Rule-based information extraction is dead! Long live rule-based information extraction systems! In *Proceedings of EMNLP*, pp. 827–832.

Cormack, G. V., Clarke, C. L., and Buettcher, S. 2009. Reciprocal Rank Fusion Outperforms Condorcet and Individual Rank Learning Methods. In *Proceedings of SIGIR*, pp. 758–759.

Demšar, J. 2006. Statistical comparisons of classifiers over multiple data sets. *Journal of Machine Learning Research*, 7:1–30.

Formal, T., Lassance, C., Piwowarski, B., and Clinchant, S. 2021. SPLADE v2: Sparse Lexical and Expansion Model for Information Retrieval. *arXiv preprint arXiv:2109.10086*.

Harris, C. 2024. BM25 vs Semantic Search: An Empirical Comparison for Domain-Specific Retrieval. *Technical Report*.

Karpukhin, V., Oğuz, B., Min, S., Lewis, P., Wu, L., Edunov, S., Chen, D., and Yih, W.-t. 2020. Dense Passage Retrieval for Open-Domain Question Answering. In *Proceedings of EMNLP*, pp. 6769–6781.

Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., and Liu, T.-Y. 2017. LightGBM: A Highly Efficient Gradient Boosting Decision Tree. In *Advances in NeurIPS*, 30:3146–3154.

Khattab, O. and Zaharia, M. 2020. ColBERT: Efficient and Effective Passage Search via Contextualized Late Interaction over BERT. In *Proceedings of SIGIR*, pp. 39–48.

Larson, S., Mahendran, A., Peper, J. J., Clarke, C., Lee, A., Hill, P., Kummerfeld, J. K., Leach, K., Laurenzano, M. A., Tang, L., and Mars, J. 2019. An Evaluation Dataset for Intent Classification and Out-of-Scope Prediction. In *Proceedings of EMNLP-IJCNLP*, pp. 1311–1316.

Lin, J., Nogueira, R., and Yates, A. 2021. Pretrained Transformers for Text Ranking: BERT and Beyond. *Synthesis Lectures on HLT*, Morgan & Claypool.

Lyu, C., et al. 2025. Attention Mechanisms for Text Classification: A Survey. *ACM Computing Surveys*.

Ratner, A., Bach, S. H., Ehrenberg, H., Fries, J., Wu, S., and Ré, C. 2017. Snorkel: Rapid Training Data Creation with Weak Supervision. In *Proceedings of VLDB*, 11(3):269–282.

Rayo, L., de la Rosa, A., and Garrido, A. 2025. A Hybrid Approach to Information Retrieval and Answer Generation for Regulatory Texts. In *Proceedings of COLING 2025*.

Reimers, N. and Gurevych, I. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of EMNLP-IJCNLP*, pp. 3982–3992.

Robertson, S. E., Walker, S., Jones, S., Hancock-Beaulieu, M., and Gatford, M. 1995. Okapi at TREC-3. In *Proceedings of TREC*, pp. 109–126.

Souza, F., Nogueira, R., and Lotufo, R. 2020. BERTimbau: Pretrained BERT Models for Brazilian Portuguese. In *Proceedings of BRACIS*, pp. 403–417.

Tunstall, L., Reimers, N., Jo, U. E. S., Bates, L., Korat, D., Wasserblat, M., and Pereg, O. 2022. Efficient Few-Shot Learning Without Prompts. *arXiv preprint arXiv:2209.11055*.

Zhang, Z., et al. 2021. Multi-Turn Intent Classification with Hierarchical Attention. In *Proceedings of NAACL*, pp. 1845–1855.